

Chapter 9: Main Memory

Memory - Ch 9-10

Address Spaces

Logical Address: Seen by CPU (Virtual)

Physical Address: Address in RAM

MMU: Hardware to map logical \rightarrow physical

(In)Memory Allocation

Base & limit Registers

Base: Start of allocated memory

limit: Source check \rightarrow if violates \rightarrow trap

Allocation strategies:

First Fit: First hole large enough

Best Fit: smallest hole big enough

Worst Fit: largest hole available

Fragmentation:

External: scattered free memory

Internal: unused memory within allocated block

Compaction: Merge free holes \rightarrow one large hole

Only possible with dynamic relocation

Paging:

Frame: Fixed-size block in physical mem.

Page: Fixed-size block in logical memory

Page table maps pages to frames

TLB: cache their stores (page # \rightarrow frame#) mapping

TLB hit: faster access, (no mem ref) TLB miss: mem

(mem. then access page in memory, TLB miss + mem access)

(translation lookaside buffer)

Logical Address (Virtual Address)

Issued by CPU; split into: Page#(p) | Offset(d)

Physical Address: needs to convert: Frame#(f) | Offset(d)

Produced by MMU: mem location

Fragmentation: words aligned

introduction in last range (avg 16-18% for Mem)

Shared pages: p/c: read-only (non-modifying)

code can be shared (e.g. library, etc.)

Chapter 10 (Virtual Memory)

Demand Paging

Only load pages when referenced

Page faults occur to non-resident page

\rightarrow trap to OS \rightarrow bring from disk

benefit: reduces main usage, # of processes

longer faults \rightarrow not load all at once \rightarrow better program

valid: invalid bit in page table setup

v: page in mem, t: not in mem, m: fault

Page fault steps (long cycle)

Page replacement

Method when v: fix frame

Algorithm

Description

Notes

FIFO

Evict oldest page

May suffer from Belady's Anomaly

Optimal

Replace page used farthest in future

Theoretical best

LRU

Replace least recently used page

Good performance; needs timestamp or stack

Clock / Second Chance

FIFO with use bit

If use bit = 1, clear and skip

Enhanced Second Chance

Uses (ref, mod)

Better approximation

Locality & Working Set

Modifying (dirty) 0/1 \rightarrow only write dirty pages to disk

Locality & Reference

temporal locality: re-use of recently accessed pages

spatial locality: reading pages likely accessed soon

Working set model

WSS(t, Δ) = pages accessed in last Δ time units

WSS(t) = size of set of pages i

D = 2 * WSS: total demand

IR: D > m (available frames) \rightarrow THRESHOLDING

Inverse Page Table

1: Entry per frame, not per process

Entry: (P#(page), frame #)

Stores mem. not stored (uses hardware)

Binding Times

Compile time: absolute address embedded

Load time: if relocatable \rightarrow absolute

Execution time: Uses MMU to bind dynamically

Segmentation

Program = logical segments (code, data, stack)

Logical Address = (segment#, offset)

Segment Table = base + limit per segment

Advantages: Logical program structure preserved

Supports sharing, protection

Disadvantages: External Fragmentation.

Effective Access Time (EAT)

$$EAT = d(TLB + M) + (1-a)(TLB + 2M)$$

d = TLB hit ratio

M = Memory access time

TLB = TLB lookup time

Page Table Structures

Hierarchical: Break large page into multiple levels

Flat: Suitable for large address spaces.

Inverted: Loading per frame, uses hashing of (P#(p), offset)

Page Table Registers: PTB(p): points to base offset

PTB(p): P#(p) in register (access validity)

Swapping: Swap process out to disk, back in when needed.

Requires dynamic relocation.

Not used in Mobile OS.

Ex: Hierarchical

IS 32(8B)

2-level paging 10 bits for page dir (NBB) Similar hierarchical 32/10

10 bits for page 12 bits

4-bit support.

EAT (Effective Access Time)

$$EAT = (1-p) \cdot \text{main access} + p \cdot (\text{page fault time})$$

Example: Mem access: 200 ns, pg fault time: 8 ms, p = 0.01

Ex: 200 + 0.01 * 8,000,000 = 8,200 ns = 0.4 ms slowdown

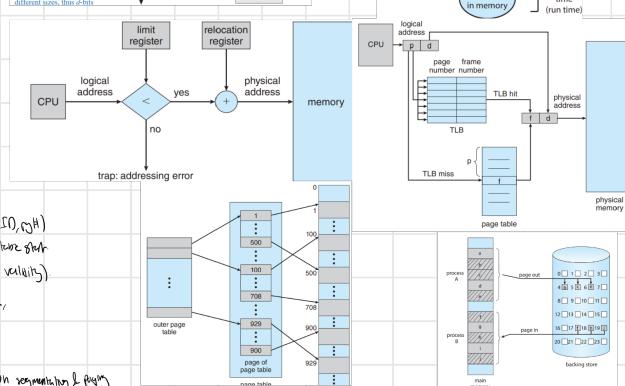
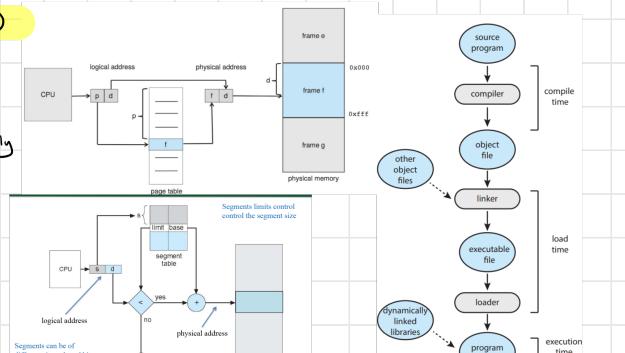
(copy-on-Write (CoW), used by fold)

Global Replacement: Local Replacement

Victim frame selected from faulting process

Victim can come from any process

Higher-priority processes get more frames



Thrashing (Too much pg. swapping)

High page fault rate \rightarrow low CPU use

System uses more processes, workers have

Solution: NBB model:

monitor 2 separate processes if $p > m$

page fault frequency (PFF) control

active memory based on fault rate

flexible degree of multiprogramming

Memory Mapping File: principle

Use mmap() to file virtual memory

File access \rightarrow memory reads & writes

Benefits: simple, VFS, interprocess communication, on-demand loading

mmap(): can map file

Kernel Memory Allocation

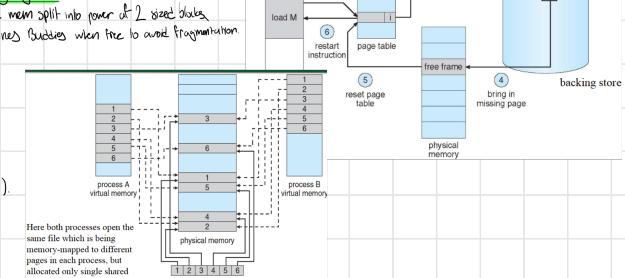
Different from user processes paging

Must allow quantity & file dictatorial

Buddy System

Free mem split into power of 2 sized blocks

combines blocks when free to avoid fragmentation



Chapter 13 - File System Interface

File concept: File is a named collection of info stored on disk

- OS converts physical storage onto logical units (files)

- Files persist across program executions

File attributes:

- Name: human readable identifier

- Type - by extension or metadata

- Location: disk address

- Size - current file size

- Protection: read/write/executable/etc.

- Timestamps: created, modified, accessed

- Owner/group - for access control

File Operations:

create(), open(), read(), write(), seek(), close(), delete(), truncate()

Open File table: holds open files into pointer (symbolic)

See: repositions file pointer: Open/close: sync between memory & disk

Direcory Structures

Type	Description
Single-level	One flat directory

Two-Level	Separate directories per user
-----------	-------------------------------

Tree	Hierarchical (common)
------	-----------------------

Acyclic Graph	Allows shared sub dirs
---------------	------------------------

General Graph	Supports cycles (w/ issues)
---------------	-----------------------------

Disk operations:

mkfs(), rmfst(), mkt(), umkt(), unmkf(), renmkf(), open(), close()

File Sharing

- File sharing across users or systems (NFS)

- Sharing access control & protection

- Shared directories & files need permissions & possibly remote access support

File System: (Ch 13-15)

File Access Methods

Method	Description	Use Case
Sequential	Read/write in order	Reading logs
Direct	Access by block #	Databases
Indexed	Uses index for quick look-up	OBs, Catalogs

File Structure

No structure: byte stream

Accord: fixed/variable size

(Complex: sometimes files executable),

File parser must understand format.

File Links:

Hard link: multiple names → same inode

Soft link (symlink): points to another path

Deletion: File removed when all hard links gone

File Protection

read, write, execute, append, delete, lock, change, attr.

Protection Models

Model

Description

UNIX permissions: Owner/group/others + rwx bits

ACLs: Fine-grained list of who can do what

Access Matrix: List of Subusers/objects = rights

Mounting & Partitions

- Mount: Attach a FS to mount point

- Unmounted FS: Inaccessible

- Each volume has its own directory tree, mounted into the global FS

File System Structure

File System resides on secondary storage (HDD/SSD). Operations are layered:

- Logical FS → directories, metadata
- File organization module → block mapping

(Basic FB) → 15-bit block # → command

I/O controller → device commands

Drivers → devices

Interface of hardware

Blocks & Tables

Boot Control Block: contains sys.info needed to boot OS from disk

File Control Block (FCB): stores structure of pic into file (for process-specific info)

Volume control block: contains all block info (total #, # of free, block size, free block ptrs)

Allocation Method: ext4 - contiguous block of 128MB

Best Fit

Fast exec, suffers from fragmentation

Linked List: pointers between blocks

Segmented: table of pointers

Indexed: index block w/ block pointers

UNIX mode: 12 direct, single, double, triple indirect

Supports huge files (4TB)

Performance Considerations

Sequential vs Random

Impacts allocation strategy

Use clustering to reduce disk seeks

Pre-read, write-behind & caching

Improve throughput

Optimize CPU vs. disk tradeoffs

Considering clusters: scan or read underlying hardware movement

Backups: allow recovery from corruption or loss

Journaling: log structure FS: track updates of appends in pur

Chapter 14 - File System Implementation

On-disk in-memory structures

Structure	Purpose
Boot (initial Block)	Boot OS from disk
Volume Control Block	Disk layout info: block count, free blocks, FCB pointers

Directory Structure	File names → index of metadata
File Control Block	Permissions, size, timestamp points to data blocks

Open File tables	Info on opened files (open wide & open pieces)
Buffers/Caches	Store data in transit between memory & disk

Blocks & Tables	Mount table: Stores mounts, MFS, FS types
Boot Control Block: contains sys.info needed to boot OS from disk volume.	System operation table: contains FCB (copy, exec, etc.)

File Control Block (FCB)	Storage structure of pic into file (for process-specific info)
Volume control block: contains all block info (total #, # of free, block size, free block ptrs)	Mount table: contains pointers to relevant entries of system-wide operation table

Allocation Method

ext4 - contiguous block of 128MB

Type	Details
Contiguous	Start/Length
Linked	Pointers between blocks
FAT	Table of pointers
Indexed	Index block w/ block pointers
UNIX mode	12 direct, single, double, triple indirect

Free Space Management

Method	Description
Bit Map	1 bit per block (0=use, 1=free), fast search

Linked List: blocks point to next free or slow search

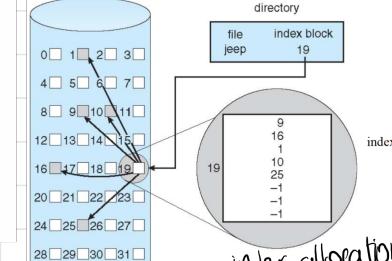
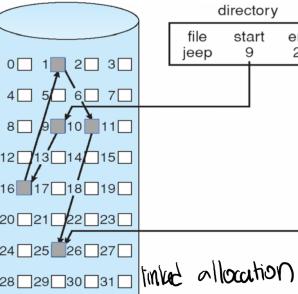
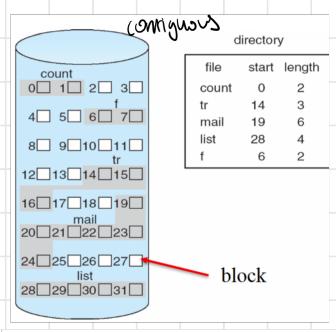
Counting: first block, stores N free addresses, down to next used

Quadratic: tracks free blocks, count of free sequential blocks

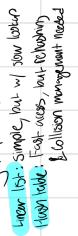
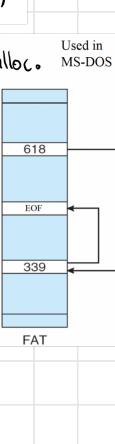
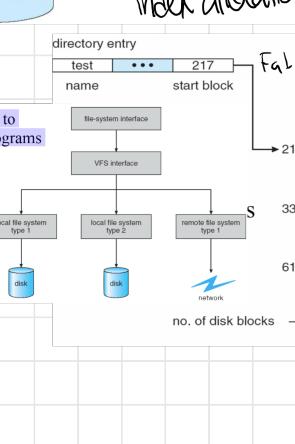
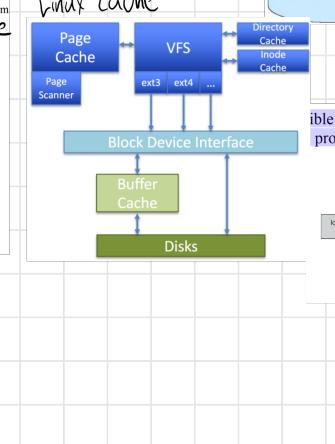
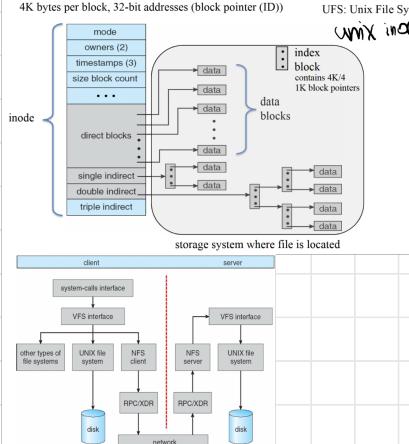
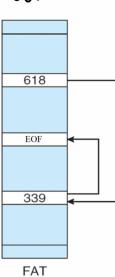
Space Map (FS): log disk activity per inodes, memory & scalability

TRIM - File system tells SSD you can erase these now.

It bridges the gap between FS free space and SSD block reuse to optimize speed and durability.



Used in MS-DOS



Chapter 15 - File System Internals

File System Basics

- A file system allows data to be stored, and retrieved
- A computer may support multiple FS types (e.g., tmpfs, ZFS, etc.)
- Storage hierarchy: physical devices → partitions → volumes → FS

File Systems Mounts

- File systems must be mounted before use.
- Mount point: a directory where a FS is attached (e.g., /home)
- Must contain valid FS structure; often validated at boot
- Some OSes (e.g., macOS) auto-mount detected volumes
- Mount Examples: mount /dev/sda2 /mnt/home
- Mounting hides existing contents of the directory

Partitions

A partition is a logical division of a disk.

Partitions can be:

- Raw: no FS (e.g., swap) Root partition: contains OS
 - Logical: contains FS
 - Volume managers can span FS across multiple partitions/disks.
- Virtual File Systems (VFS)** uses inodes; holds index or network file details
- Provides abstraction layer for all FS types.
 - Defines a common API for creating files.
 - Each FS provides its own VFS operating (e.g., lookup, read).
 - Used by UNIX, LINUX, Windows (Win32 API under VFS model)

Memory File Systems

Systems use
MFS (Memory File System) UNIX-based MFS-driven

CIFS/SMB Microsoft, authentication via Active directory

FTP/Web Simple transfers, less integration

NFS Protocol

- Stateless: each request contains all needed info
- Uses Mount Protocol & Path-name Translation

Consistency Semantics

Model Description

UNIX Changes visible immediately

Session Semantics Variables only visible after file closed

Immutable-Shared No changes once file is opened

OS must coordinate how file changes are seen by different process

File Semantics in Practice

- Systems like Linux use virtual FS that compute file contents dynamically
- Immutable FS useful for concurrency (e.g., copy-on-snapshots)

File Sharing & Partition

- Multi-user system: need controlled sharing
- Metadata includes

• Owner • Permissions

File sharing handled via:

• local mount

• remote access (e.g., NFS, CIFS/SMB)

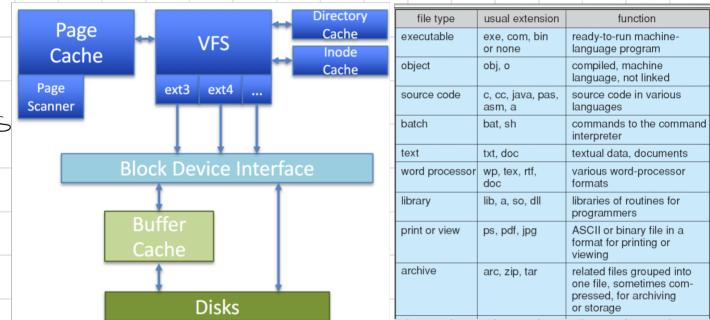
(Case Study) NFS (Network File System)

Mount protocol authenticates & mounts remote FS

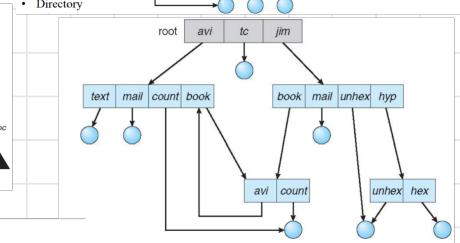
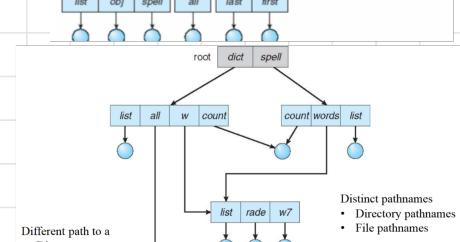
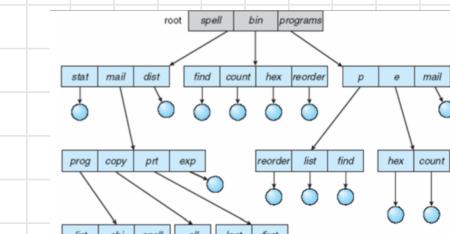
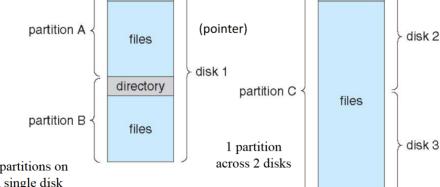
NFS Protocol

- Starts communication
- Handles operations like read, write, lookup, setattr

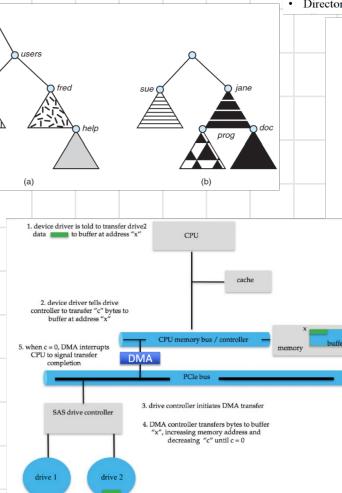
May suffer staleness: issues unless careful caching is used



file type	usual extension	function
executable	exe, com, bin or none	ready-to-run machine-language program
object	obj, o	compiled, machine language, not linked
source code	c, cc, java, pas, asm, a	source code in various languages
batch	bat, sh	commands to the command interpreter
text	txt, doc	textual data, documents
word processor	wp, tex, rtf, doc	various word-processor formats
library	lib, a, so, dll	libraries of routines for programmers
print or view	ps, pdf, jpg	ASCII or binary file in a format for printing or viewing
archive	arc, zip, tar	related files grouped into one file, sometimes compressed, for archiving or storage



- An access matrix is one way to represent policy.
 - Frequently used mechanism for describing policy
 - Columns are objects
 - Subjects are rows
 - Cells are rights
- To determine if S_i has right to access object O_j , find the appropriate entry.
- Succinct descriptor for $O(|S| * |O|)$ entries
- There is a matrix for each right



O ₁	O ₂	O ₃	
S ₁	Y	Y	N
S ₂	N	Y	N
S ₃	N	Y	Y

Chapter 12: I/O Systems

I/O System Overview

- I/O is slow relative to CPU/memory → needs OS optimization
- OS abstracts wide variety of devices: Device drivers: encapsulate
 - Character/block, sequential/random access, device details.
 - Shared/dedicated, read-only/read-write.

I/O Hardware Components

Component	Role
Device Controller	Manages a device (Registers + Buffer)
Bus	Transfers data/commands between CPU, mem, I/O
Port	Interface to communicate w/ controller
Fiber Channel (FC)	Complex controller (Separate circuit board)
Memory mapped I/O	device registers mapped into sys. mem space
BUS: device char or shared direct access:	
PCI, PCIe	Expansion bus
SCSI (SATA)	Serial ATA
IDE	Parallel ATA
USB	Universal Serial Bus

I/O Communication Methods

Method	Description	Pros	Cons
Polling	CPU checks status repeatedly	Simple	CPU waste
Interrupts	Device signals CPU on readiness	Efficient	Overhead, needs handler
DMA	Device writes data directly to memory	CPU offload	Set up complexity

DMA: Direct memory access → best for bulk transfers (disk, NPC)

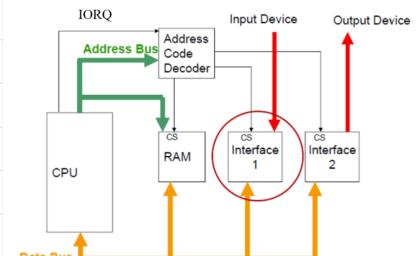
DMA steps: CPU sets up → DMA controller handles block → interruption complete

Application I/O Interface

- Provided via system calls (read(), write())
- I/O Modes
 - Blocking: waits until I/O completes
 - Non-Blocking: returns immediately (poll status later)
 - Asynchronous: completion notified via signal/call back
 - Vectorized I/O: multiple buffers in one I/O operation

Devices vary in many dimensions:

- Character stream **or** block
- Sequential **or** random access
- Synchronous **or** asynchronous (block?)
- Shareable **or** dedicated
- Speed of operation
- Read/write, read only, write only.

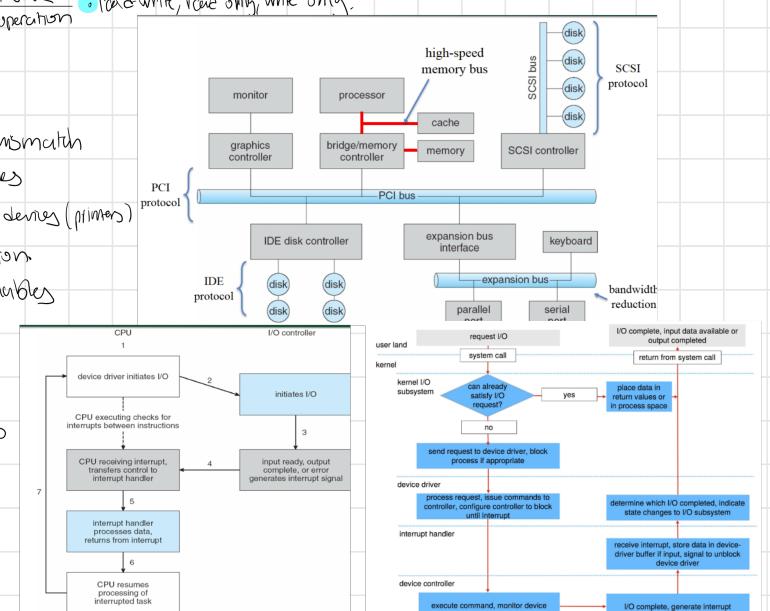


Kernel I/O Subsystem

Responsibilities:

- I/O scheduling: optimize request order
- Buffering: temp storage to smooth mismatch
- Caching: store frequently accessed blocks
- Spooling: queue jobs for non-shareable devices (printers)
- Error handling, protection, device reservation

Buffering decouples producer/consumer, enables prefetching & delay tolerance



I/O Request Life Cycle:

- App issues sys call (read())
- OS checks cache → if miss, schedule device I/O
- DMA handles transfer (if applicable)
- Device triggers interrupt on completion
- OS returns app w/ data/result

I/O Scheduling

- Similar to CPU, but for I/O queuing (esp. disks)
- Prioritize time-sensitive ops (e.g. virtual memory in)
- Batch reads/writes to reduce overhead.

Device Drivers

- OS component knows how to interact w/ a specific device
- Each driver implements its own interface
- Layered Model: **Top layer**: device-independent code
Bottom layer: hardware-specific logic

Streams (Advanced UNIX I/O)

- Message passing architecture for I/O pipelines (System V)
- Enables modular drivers. Not common today, though.

Performance Optimization

Technique	purpose
Buffering	Smooth I/O rotation
Read-ahead	Prefetch likely needed data
Write-behind	Delay writes for batching
Double buffering	Avoid stalls by waiting
Interrupt coalescing	Minimize interrupt frequency

CS (c)